

Mode of Engagement: A Prior Classification Layer

Part of `foundations/` — see `../brief.md`. **Status: PENDING, but now first in the working order (resequenced 2026-09-19)** — the competitive-practice mode (§3) has a full, evidenced definition as of

`../posts/post_2_unit_of_analysis/` (this project's "Post 2"), which is what actually resolved the sequencing question §6/§8-task-6 originally left open on 2026-09-17. Still PENDING rather than ADOPTED: the narrative-consumption and social-creative modes remain undeveloped (§3), and §5's cross-mode comparability work hasn't started.

Post 2's own scope narrowed as of its Draft 4 (2026-09-19): the definition below is unchanged, but the post now stops there rather than continuing into "how do we compare and contrast esports" (its own "Problem 2") — that empirical work (genre/niche porousness, creator crossover, region/language confounds) has been moved to a separate, not-yet-titled follow-up post. `../notebooks/creator_crossover_subset.ipynb` and `../notebooks/region_language_confound.ipynb`, built for that section, are parked in `../posts/post_2_unit_of_analysis/` pending that post's own folder — real, correct evidence, just not evidence *for this post* anymore. See §9's last bullet.

1. Why this exists

Raised directly 2026-09-17, out of a felt problem with `esports_vs_game_fandom_signal_model` rather than an abstract exercise: strip out the narrow slice of people who watch a title's esports scene but never play it, and "esports fandom" looks inseparable from "game fandom" generally — esports viewership reads as a behaviour that emerges in people who already play, not a distinct audience. (Testable, not yet tested — see §5, task 4.)

But the reasonable-sounding fallback, "fine, then just compare against *game* fandom broadly," doesn't survive scrutiny either. Is God of War fandom the same kind of thing, at a sociological level, as VALORANT fandom? One resembles being a film fan; the other resembles playing in a 5-a-side league. **"Game" (short for "video game") names a delivery substrate — the fact that the activity is software — not a mode of engagement or the role the activity plays in someone's life.** Treating "game fandom" as a coherent category because its members share a technology is the same move as treating "cinema fandom" as one thing because action films and documentaries share celluloid.

2. The core claim

The unit of analysis this project needs isn't "video game" (a technology category) or a single esports/game-fandom split — it's *mode of engagement*: the behaviour-group a title's fandom actually belongs to, defined by what the activity is for in people's lives, not by what it runs on. Mode of engagement sits *above* genre, platform, and region in the classification stack — those describe a title; mode describes what kind of relationship people have with it.

This is deliberately fluid, not a clean partition — see §4. The point of naming it isn't to force every title into exactly one box, but to stop silently assuming two titles' fandoms are comparable just because they share a genre tag or a delivery technology.

3. Candidate modes (draft, not final)

Three, provisionally, drawn from the titles already in this project's own scope:

- **Competitive-practice.** Now has a full, evidenced definition, not just the sketch below — see `../posts/post_2_unit_of_analysis/` (drafted 2026-09-17 to 2026-09-19): **any persistent, competitive,**

community practice of play, mediated by software. No professionalization gate — casual and competitive players differ by degree, not by type, and watching/playing are treated as the same underlying practice rather than two audiences to disambiguate. Deliberately not scoped to "video game": chess, poker, the Excel and GeoGuessr world championships all qualify on the same terms, which is the whole point (§3's chess/poker note below anticipated this before the post resolved it). VALORANT, Counter-Strike, Dota 2, Street Fighter all qualify straightforwardly. `esports_vs_game_fandom_signal_model`'s `broadcast_tier`-based work is now understood as *detecting* this mode in this project's Twitch data specifically, not as a separate, competing definition of it (§6). **Draft 4 sharpens the watch/play claim further, not just "differ by degree" but an explicit list of what the practice actually consists of:** playing, watching, Discord/chat participation, social-media engagement with pros, event travel, merchandise purchase — "a soup of behaviours," not one activity. That list is a first, prose-form draft of §5's still-undecided cross-mode-comparable component dimensions (see §8 task 3) — a starting point to formalize, not yet a measurement framework.

- **Narrative-consumption.** Closer to a film or book than a sport — authored story, largely solitary or para-social engagement, no ranked or competitive structure. God of War is the reference case; none of the 23 currently-tracked titles are a clean example, since this project's tracked list was built esports-first (see `../brief.md`).
- **Social/creative-sandbox.** Community forms around a shared space or shared activity, not around competition or an authored story. GTA Online's NoPixel roleplay scene is the case already sitting in this project's own data (`platform_viewership_snapshots.content_segment`, `research/wider_game_fandoms/gta_launch_case_study/`) — intense, sustained fandom with no ranked ladder and no authored narrative arc driving it.

These are not substrate-exclusive. Chess and poker are not video games and predate digital delivery entirely, yet both show the same competitive-practice fandom shape this project associates with esports — ranked play, professional circuits, livestreamed tournaments, a watch-but-play-casually-or-not-at-all audience. That's evidence "digital" is infrastructure, not the sociologically load-bearing part of the category: what matters is a legible competitive structure plus a broadcast mechanism, and most of this project's titles happen to deliver both through video-game software because that's the infrastructure that existed when they formed (a direct connection to `../technological_contingency_of_fandom_formation/`'s own claim — worth making explicit there once this model firms up).

4. Fluidity: mode is per-context, not always per-title

A single title can occupy more than one mode depending on which part of its audience or which mode of play is in view. GTA V's mainline campaign is narrative-consumption; GTA Online's roleplay servers are social-creative sandbox; neither is competitive-practice. Treating "GTA V" as one title with one mode would erase exactly the distinction this model exists to draw. **Mode may need to be a property of a community/context (a subreddit, a stream category, a server type) rather than a title-level tag in `config/titles.yaml`** — an open question, not resolved here (§5, task 1).

5. Why comparability matters, and what it requires

Two things need to be true at once, per direct instruction 2026-09-17: mode-of-engagement should let this project **compare fandom components across modes** (is a competitive-practice fandom's community intensity comparable to a social-creative sandbox's, on some shared axis?) as well as **compare subtypes within a mode** (VALORANT vs. Counter-Strike, both competitive-practice, as `fandom_and_niche_model` already does). The second is what this project already does. The first requires something that doesn't exist yet: **fandom decomposed into component dimensions that aren't mode-specific**, e.g. ritual/routine engagement frequency, creator/parasocial attachment, competitive-identity investment, social-coordination intensity, narrative/lore investment, collecting or completionism behaviour. A competitive-practice fandom and a social-creative one will score differently across those dimensions, but the dimensions themselves should be

measurable in both — that's what makes "compare across modes" a real claim rather than a category error repeated one level down. Not designed yet (§8) — though see §7 for a first prose-form candidate list Post 2's Draft 4 already sketched.

6. Relationship to the other three foundational pieces

- **esports_vs_game_fandom_signal_model** — **reconciled 2026-09-19, not left ambiguous**. Its old operational definition (§4 of that brief: esports fandom = viewership in `broadcast_tier ∈ {primary_official, detected_costream}`) is retained, but reframed as a *measurement operationalization* of this model's competitive-practice mode, not a competing conceptual definition of esports fandom itself. That brief's §2 is rewritten to drop the game-fandom-vs-esports-fandom disambiguation project entirely and cite this brief instead — see that brief's own §2 for the rewrite.
- **fandom_and_niche_model** — **reconciled 2026-09-19**. Its old §2 scoping definitions ("Esports": professional/cash-gated; "Successful esports": a 2-year A-tier history as the entry bar) directly conflicted with this model's definition, which drops professionalization entirely. Fixed by citing this brief for what counts as an esports at all, and reframing "Successful esports" as a maturity/scale threshold measured *within* the competitive-practice-mode category rather than the gate into it — `get_success_milestone`'s pipeline is unaffected, only what its output is understood to mean. This was also the concrete argument for sequencing mode *before* niche: a `{genre × platform × region}` cell built on an unrevised "Esports" definition would have been comparing across the exact category gap this model exists to name.
- **technological_contingency_of_fandom_formation** — see §3's chess/ poker point: if competitive-practice fandom predates digital delivery, then "digital substrate" is one technological condition among several that can produce this mode, which is squarely that brief's own claim (fandom formation is contingent on the tech/social conditions present when it formed) rather than a separate idea living here by accident. **A second, opposite-direction connection surfaced in Post 2's Draft 4 postscript**: citing a 2023 Deloitte report on traditional-sports fans increasingly engaging via second-screen, social platforms, betting, and VR interest — "are there any sports that aren't esports now?" That's the same contingency claim running forward instead of backward: not "esports emerged because specific tech conditions existed," but "traditional sports are being reshaped toward competitive-practice-mode behaviour as those same conditions arrive for them, later." A genuine candidate second case study for that brief, not yet written up there — currently just an observation in a postscript.

7. A practical implication, from Post 2's Draft 4: attention over video games

Not a research question, a business-facing consequence worth recording since it's new to this brief and speaks directly to the audience `../brief.md` names (esports-industry professionals): **if the video game is an interchangeable substrate and the human network is the actual unit of analysis, then the thing that generates value is attention, and attention is a property of the network, not of any one piece of software**. Practically, that argues against scanning multiplayer-title release calendars as the primary way to find new esports-adjacent opportunity — the Excel/GeoGuessr World Championship examples (§3) are the concrete illustration — and toward looking for *any* substrate capable of generating a persistent, competitive, community practice of play, regardless of whether it was built as a game at all. Not yet connected to anything measurable in this project's own data; recorded here as a live idea Post 2 raised, not a task with a method behind it yet.

8. Task list

1. Decide whether mode is title-level or context-level (§4), and if context-level, what the actual unit is (subreddit, stream category, server/game-mode) — needed before any title in `config/titles.yaml` can be tagged. Still open.

2. ~Draft structural markers for the competitive-practice mode specifically~ — **done, superseded by a cleaner result.** `../posts/post_2_unit_of_analysis/` didn't produce a checklist of structural markers, it produced an actual definition (§3) that doesn't need one: persistent + competitive + community + software-mediated, no professionalization gate. Whether this reproduces or disrupts the current genre taxonomy if applied to all 23 tracked titles is still open — that's now task 2a below. 2a. Apply the §3 definition to all 23 tracked titles and check whether it reproduces or reshuffles the current genre-based grouping. Not started.
3. Design the cross-mode-comparable fandom-component dimensions in §5. Not started; this is the harder and more novel half of the work here.
4. Test the premise in §1 directly: does Twitch viewership decouple from Steam playerbase for competitive-practice titles (viewership sustaining or growing while playerbase flattens), which would be evidence of a real watch-only segment, against a stable/correlated relationship, which would support the "viewing is emergent from playing" reading. Both series already exist in this project's data. Still not run — the most concrete, cheapest-to-answer open task here.
5. ~Reconcile with `esports_vs_game_fandom_signal_model` ~ — **done, 2026-09-19. See §6.**
6. ~Decide sequencing~ — **done, 2026-09-19.** This brief is now position 1 in `../brief.md`'s working order.

9. Open items

- Whether three modes (§3) is the right cardinality, or an artifact of this project's own esports-first title list. `wider_game_fandoms` resuming may surface titles that don't fit cleanly (a live-service game with light PvP and heavy narrative content, for instance).
- The chess/poker evidence (§3) is a suggestive analogy, not a tested claim — this project has no data on either, so it can motivate the model but shouldn't be cited as validating evidence.
- Whether "mode" and "niche" end up as two independent axes (`{mode × genre × platform × region}`) or mode subsumes/replaces one of the existing niche-cell dimensions once §8 task 2 is done.
- This brief was written directly from a research-journal reflection (`docs/research_journal.md` , 2026-09-17 entry) and then from `posts/post_2_unit_of_analysis/`'s own drafted argument, rather than from a notebook finding — unlike the other three foundational pieces, the competitive-practice mode's definition (§3) is argued, not yet measured. Task 4 above is the actual empirical test that definition is still waiting on.
- `../notebooks/creator_crossover_subset.ipynb` and `../notebooks/region_language_confound.ipynb` are now orphaned from the post they were built for. Both are evidence for a related but distinct claim (genre bins don't track real creator affinity; region/ language is a real but non-general confound on that pattern) — real data, not a direct test of this brief's own central definition, and built for Post 2's "Problem 2" section before that section moved to a separate, not-yet-titled follow-up post (2026-09-19, Draft 4). Left parked in `../posts/post_2_unit_of_analysis/` pending that post getting its own folder — a housekeeping item, not a research one.

10. A candidate sequence for §5's component dimensions (2026-09-25) — and why `milestone_year` isn't a network-maturity measure

Raised directly out of the `post_3_anatomy_of_a_network` `milestone_year` robustness work, not designed from scratch here: `get_success_milestone()` (this project's only "when did a title's scene mature" measure, load-bearing across `../posts/post_3_anatomy_of_a_network/brief.md` , `../esports_lifecycle_and_maturity/brief.md` , and `docs/milestone_reconciliation.md`) is confirmed, from the actual code, to be entirely a professionalized-competitive-infrastructure signal: the last year of the earliest 2-consecutive-year window containing an S/A-tier (Liquipedia's own top two tiers) tournament each year, spanning 2+ continents cumulatively. Viewership is evaluated

(`viewership_check`) but non-gating and frequently `None` . No signal about grassroots play, casual playerbase, audience size, social behaviour, event attendance, or purchasing enters the calculation anywhere.

That's narrower than what this whole project has been using it to mean. Every notebook built on "milestone_year measures when the community/ network matured" — post_3's era hypothesis chief among them — has actually been testing "when the *professionalized competitive tier* reached a specific, fairly high, international bar." A real thing, but one specific rung of something bigger, not the network's maturity as a whole.

The sharper diagnosis: this makes milestone_year a *lagging* indicator of network maturity, not a leading one. A professional/elite competitive tier requires an existing playerbase large enough to produce elite players and organizations willing to invest in them — it cannot exist before a real base of play does. Wherever milestone_year has been read as "the moment this title's community formed" (its use throughout this project so far), the more accurate reading is "the moment the community had already matured enough, for long enough, to support professionalization" — a downstream consequence of network maturity, not the event itself. This also directly explains why the independent changepoint-detection check in post_3's brief lands at or after milestone_year for every title it could test (0 to +3 years, never negative): Twitch viewership — a much closer proxy for Watch — isn't tracking milestone_year, it's usually leading it slightly.

A candidate structure for what actually needs measuring, proposed directly and not yet tested:

Play → Elite Play → Professionalised Play → Watch → Socialise → Attend → Purchase

— not strictly linear. Watch is plausibly *ahead of* Professionalised Play, not behind it, for any title released after livestreaming technology matured: a title can build a large spectator audience around high-level-but-not-yet-institutionally-professionalized play (arguably closer to what actually happened for this project's own deepest case study, VALORANT). Socialise/Attend/Purchase likely have no single fixed position relative to each other or to Watch — a critical-path model probably breaks down past Watch, and a second, parallel fork is likely needed for the *streaming/creator* side of the network (a distinct population from players, feeding back into Watch) rather than folding creators into one single-file sequence.

This is the same idea as §5's still-undisigned "cross-mode-comparable fandom-component dimensions" and §7's prose-form list from Post 2 Draft 4 (playing, watching, Discord/chat, social-media engagement with pros, event travel, merchandise) — arrived at independently, from interrogating milestone_year's limits rather than designing §5 from a blank page. Two independent routes landing on the same rough list is a good sign it's pointing at something real.

What this project can and can't currently measure of that sequence: - **Play / Elite Play** — not measured directly. No grassroots participation or matchmaking-tier data for any title (Steam player counts are the closest proxy, and they're aggregate, not skill-tiered). - **Professionalised Play** — `get_success_milestone()` 's current definition. A real measurement, just of one specific rung. - **Watch** — `viewership_snapshots` (Twitch), the Kaggle historical import, YouTube. By far the best-covered component in this project's data. - **Socialise** — no direct source. Reddit was investigated and shelved (API access blockers — see `project_creator_crossover_and_reddit_status.md`); Discord has never been pursued. - **Attend** — no source exists or is planned. Neither physical event attendance nor live-viewership-of-events-specifically is collected anywhere in this project. - **Purchase** — no source exists. Steam wishlist/purchase data is publisher-only (`reference_steam_wishlist_api.md`); no merch/ancillary- spend source has been considered.

Three of seven candidate components are entirely dark to this project's current data, and the strongest-covered one (Watch) is exactly the one milestone_year's current framing leaves out.

Build this a priori or induce it from data? Both, but sequenced — and the first test is already affordable. This project already has two of the components measured for every title (Professionalised Play via milestone_year's tournament data, Watch via Twitch/Kaggle). The cheap, immediate first test: for each of the 23 titles, does Watch activity visibly precede, coincide with, or lag Professionalised Play's own milestone_year, and is the direction consistent across the

generational communities already found via Leiden clustering? Answerable with data already in hand, no new collection needed — and it should be run before any larger framework gets built on top of an unchecked intuition.

Scope of impact, flagged rather than immediately actioned. This doesn't invalidate what `get_success_milestone()` actually computes — that computation is correct for what it measures. It complicates every place this project has `read_milestone_year` as "network/community formation," which is most of the era-hypothesis work in `../posts/post_3_anatomy_of_a_network/brief.md` and touches `../esports_lifecycle_and_maturity/brief.md` and `docs/milestone_reconciliation.md`. Not rewriting those now — recorded here as the reason a future pass through that work needs to distinguish "when did professionalization happen" from "when did the network mature," rather than treating them as the same question.

Task list addition (see §8 numbering above; appended here rather than renumbering the existing list):

7. Test whether Watch (Twitch/Kaggle viewership) leads, lags, or coincides with Professionalised Play (`milestone_year`) per title, using data already collected — the cheap, immediate first test of the sequence claim above. Not started.
8. Decide whether Socialise/Attend/Purchase are worth pursuing a data source for at all, given none currently exists and CLAUDE.md's ethical-scraping constraints likely rule out several obvious shortcuts (e.g. scraping event ticketing or merch storefronts). Not started — may resolve to "structurally out of scope" rather than a data task.

11. Behaviours and artifacts of competitive-practice networks, drafted (2026-10-03) — and a three-level structure this surfaces

Raised directly, out of the `post_3_anatomy_of_a_network` `milestone_year` discussion (§10): §10's ladder (Play -> Elite Play -> Professionalised Play -> Watch -> Socialise -> Attend -> Purchase) and §7's prose-form "soup of behaviours" are formalized here into one fixed table, per behaviour: a candidate artifact (what it leaves behind, measurably) and this project's actual current data status against it. **Scoped explicitly to the competitive-practice mode only** — not a universal behaviour taxonomy across all three candidate modes (§3). Narrative- consumption and social-creative-sandbox almost certainly need their own, differently-shaped lists (neither has an "Elite Play"/"Professionalised Play" equivalent at all — there is no ranked ladder to ascend in God of War), not yet drafted.

#	Behaviour	Status	Candidate artifact(s)	Data status in this project
1	Play	existing (§10)	Steam concurrent-player count	Have — live, 11 titles, aggregate/flow not stock (<code>collectors/steam_poll.py</code>)
			Wikipedia article pageviews	Have, unused this way — <code>collectors/wikipedia_pageviews_pull.py</code> exists
2	Elite Play	existing (§10)	Grassroots (tier-3/4) tournament volume	Have, unused this way — sits in <code>tournaments_lpdb_competitive</code> , currently only used as Pro Play's exclusion set
			Ranked/MMR distribution	Don't have — no matchmaking-tier data anywhere
3	Professionalised Play	existing (§10)	Tournament records, prize pools, <code>milestone_year</code>	Have — best-built pipeline in the project
			Roster/org stability	Have — the pro-cohort thread (<code>research/foundations/notebooks/pro_player_*.ipynb</code>)
4	Watch	existing (§10)	Twitch/YouTube viewership, Kaggle	Have — best-covered behaviour overall

#	Behaviour	Status	Candidate artifact(s)	Data status in this project
			historical	
5	Discuss	<i>proposed addition</i>	Liquipedia edit-history volume	Have, unused this way — connector exists, edit activity never pulled as its own signal
			Reddit, Discord	Don't have — Reddit shelved (API blockers), Discord never pursued
6	Socialise	existing (§10)	Reddit, Discord	Don't have — same sources as Discuss; the two verbs may be practically inseparable at the artifact level even if kept conceptually distinct
7	Attend	existing (§10)	Event scale (team/participant counts)	Partial — tournament metadata gives <code>team_number</code> , not attendance
			Ticket sales, venue capacity	Don't have
8	Purchase	existing (§10)	Wishlist/sales data	Don't have — publisher-only (confirmed, <code>reference_steam_wishlist_api.md</code>)
			Sponsorship/merch deal announcements	Don't have — not collected
9	Create	<i>proposed addition</i>	Liquipedia guide/page authorship	Have, unused this way — same connector as Discuss's artifact, different edit-type if distinguishable
			Fan-platform content (art, edits, cosplay)	Don't have — likely out of scope

Discuss and Create are proposed, not settled — flagged as open rather than folded into the existing §10 ladder silently. Discuss is a small addition (already implicit in §7's "Discord/chat participation," just not yet its own rung). Create is a larger, structurally different one: it's a plausible candidate for the "second, parallel fork... for the streaming/creator side of the network" §10 already said the main Play-through-Purchase chain would eventually need, rather than another item on the same single-file sequence — fed by Watch (people who watch become the ones who create guides/content/edits), not by Play directly.

What the table surfaces operationally: three artifacts are "have, unused this way" — Steam concurrent-players, Liquipedia edit history, and grassroots tournament volume — meaning Play, Discuss, and Elite Play all have real, already-collected data sitting idle, not requiring new collection. Larger immediately-actionable surface than assumed before this table existed.

A three-level structure this draft makes explicit, clarifying §5's open "cross-mode-comparable dimensions" problem rather than just restating it: (1) **Mode** (§3) — which kind of relationship a title/ context has; (2) **mode-specific behaviours** — this table, valid only for competitive-practice; (3) **§5's still-undesigned cross-mode- comparable dimensions** (ritual/routine frequency, parasocial attachment, competitive-identity investment, social-coordination intensity, narrative/lore investment, completionism) — not verbs, but intensity-axes abstracted enough to be computed from whatever mode-specific behaviours actually apply, so a competitive-practice fandom and a social-creative- sandbox fandom can be compared on the same axis despite sharing almost no behaviour-vocabulary. Level 3 is presumably derived *from* comparing several modes' own level-2 tables side by side — which means drafting narrative-consumption's and social-creative-sandbox's own behaviour/ artifact tables (even argued-not-measured, matching competitive- practice's own current evidentiary status) is likely a real prerequisite for designing level 3 at all, not a separate, lower-priority task.

Task list addition:

9. Decide Discuss and Create's status (fold into the existing ladder, keep as open/provisional, or reject) -- not decided here.

10. Draft narrative-consumption's and social-creative-sandbox's own level-2 behaviour/artifact tables, as a likely prerequisite for §5 task 3 (cross-mode-comparable dimensions) rather than independent of it. Not started.