

Title geography, two ways: tournament host country vs. language-mix-implied country

Purpose: answers a direct capability question (2026-09-09/10) — can this project's data be shown on a color-coded world map, per title? Two genuinely different sources, shown side by side rather than blended, since they answer different questions:

1. **Where has a title's competitive scene actually been hosted** — `tournaments.country`, a real per-event fact, no proxy involved.
2. **What countries a title's Twitch audience is implied to skew toward** — `language_mix_snapshots.language_code`, a genuine proxy (a language isn't a country), and a much weaker claim than #1.

Both maps hit the same core problem at different points, worth stating up front rather than discovering per-chart: real-world geography doesn't collapse cleanly onto "one country per data point." A tournament can be labeled "Europe" instead of a specific country; a language like English or Spanish is spoken across many countries with no single dominant one. Every country-level number below is built from a **deliberately incomplete, explicitly documented** normalization — rows that don't map confidently to one country are excluded and counted, not guessed. Treat every map as a floor on real coverage, not a full picture.


Map 1 — tournament host country

`tournaments.country` is a raw Liquipedia field, never normalized before now. Checked directly (2026-09-09/10) against every tier-1/2 tournament row: it's a real mess — actual countries in wildly inconsistent form ("USA" / "usa" / "US" / "United States" / "united states"), lowercase ISO-2 codes mixed in ("kr", "de", "jp"), and, separately, a huge share of rows that aren't a country at all — continents and macro-regions ("Europe", "CIS", "MENA", "SEA", "Nordic", "Benelux") or placeholders ("World", "TBD"). None of that was assumed — every distinct value in the column was pulled and classified by hand.

```
16300 tier-1/2 tournament rows with a country value
 8429 (51.7%) map to a specific country
 7870 (48.3%) are a continent/region/placeholder, excluded
  1 (0.0%) unmapped (genuinely unclear) -- not guessed
```

Nearly half of all tier-1/2 tournament rows use a continent or region, not a country, as their location field — a real finding about how Liquipedia editors record tournament geography, not a gap in this project's parsing. Read every map below as "of the events that *did* name a specific country," not "of every event."

Tournament host country, share of tier-1/2 events (darker = more events hosted there)



Read: Dota 2, League of Legends, and StarCraft II (not shown here, richest of all 23 titles at 52 distinct countries — checked while picking this set) have genuinely global, many-country hosting patterns. Counter-Strike and VALORANT look almost **empty** on this map — not because their scenes are geographically thin (both are among the largest tracked titles by every other metric in this project), but because their tournament organizers overwhelmingly label events by region ("Europe," "Americas," "World") rather than host city/country — confirmed directly: Counter-Strike has only 10 tier-1/2 tournament rows at all (out of 11,434 total rows for the title — the rest is B/C-tier grassroots data, not eligible for a tier-1/2-only map), and of those 10, only 3 distinct countries are ever named; the other 7 use a region label. **This map type structurally favors titles/organizers that record a specific host country** — it is not a neutral measure of geographic reach, and CS/VALORANT's near-blank maps are the clearest illustration of that in this dataset.


Map 2 — language-mix-implied country

The counterpart to `notebooks/niche_membership.ipynb`'s `LANGUAGE_TO_REGION` (continent-level) — extended one level further to specific countries, **more conservatively than that mapping**, not just a finer version of it. Two languages that were "safe" at continent granularity there (`zh`, `ar`) are excluded here entirely: Chinese spans mainland China, Taiwan, Hong Kong, and Singapore, and Twitch is itself effectively inaccessible in mainland China, so a `zh` viewer skews toward Taiwan/HK/diaspora, not a single confident country; Arabic spans dozens of

countries with no dominant one. `en / es / other` stay excluded for the same reason the original notebook excluded them.

Counter-Strike 2	73.3% of viewer-time maps to a specific country
Dota 2	77.1% of viewer-time maps to a specific country
League of Legends	40.9% of viewer-time maps to a specific country
VALORANT	36.7% of viewer-time maps to a specific country
Mobile Legends: Bang Bang	93.3% of viewer-time maps to a specific country
Free Fire	29.7% of viewer-time maps to a specific country

Language-mix-IMPLIED country, share of mappable viewer-time (darker = more implied viewership)



Mobile Legends: Bang Bang's map was the one surprise, and a deeper dig (2026-09-10) found the first-pass explanation here was overstated. The map shows overwhelming Russia concentration (`ru` at **89.9%** of all its viewer-time this window). The first version of this note attributed that mostly to one live CIS-region tournament overlapping the capture window — checked more rigorously since, and that explanation covers only a small part of it: only **10.5%** of MLBB's Russian viewer-time is tournament-broadcast-tagged (`broadcast_tier='detected_costream'`); the other ~90% is ordinary, untagged community streaming, spread across **422 distinct Russian-language channels — 53.4% of MLBB's entire 791-channel Twitch creator base**, with no single channel dominating (top 10 channels are only 50.9% of the Russian total) and a daily share that stayed flat at 89.6%-92.1% across all 10 days of the window, not spiking around specific match dates. That's a broad, sustained community signal, not a tournament artifact. Full investigation, including a cross-title check (PUBG Mobile shows an

even stronger, **fully organic** version of this pattern) and absolute-scale numbers, not just share: `notebooks/mlbb_russia_deep_dive.ipynb`.

Caveats

- **These two maps measure genuinely different things and should never be read as agreeing or disagreeing with each other without checking their own coverage numbers first.** A title can look "hosted everywhere" on Map 1 while its audience-language map is thin, or vice versa — that's real information, not a contradiction to resolve.
- **Every country shown is a floor, not a ceiling**, on both maps — rows that don't map confidently are excluded and counted, never guessed. Re-derive `COUNTRY_TO_ISO3` / `LANGUAGE_TO_COUNTRY` if either project's underlying data changes shape enough that the exclusion-rate numbers above look meaningfully different.
- **language_mix_snapshots is still a short window** (`notebooks/niche_membership.ipynb` 's own caveat applies here identically — a few days of hourly polls, not a seasonally-averaged baseline, and exposed to the same live-tournament contamination risk that notebook checks for directly). Re-run this notebook once that table has weeks or months of history before treating any single country's share as settled.
- **Not every one of the 23 tracked titles is shown** — `TITLES` above is a deliberately chosen 6-title comparison set, not a claim that the other 17 don't have interesting maps. Swap `TITLES` and re-run for any other title; the two normalization dicts above already cover every language/country actually seen across all 23.